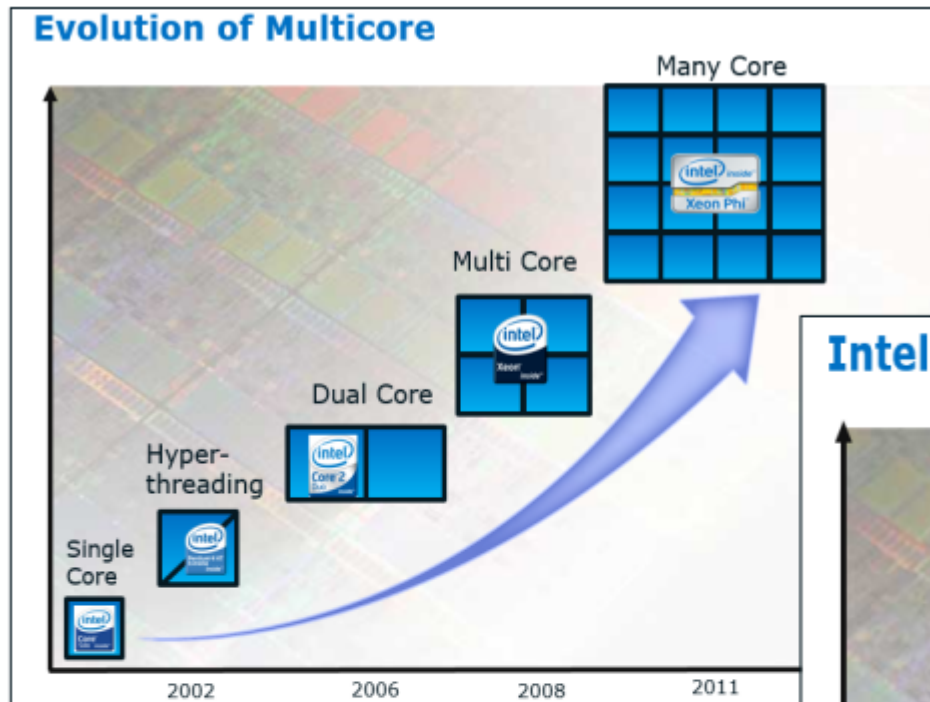


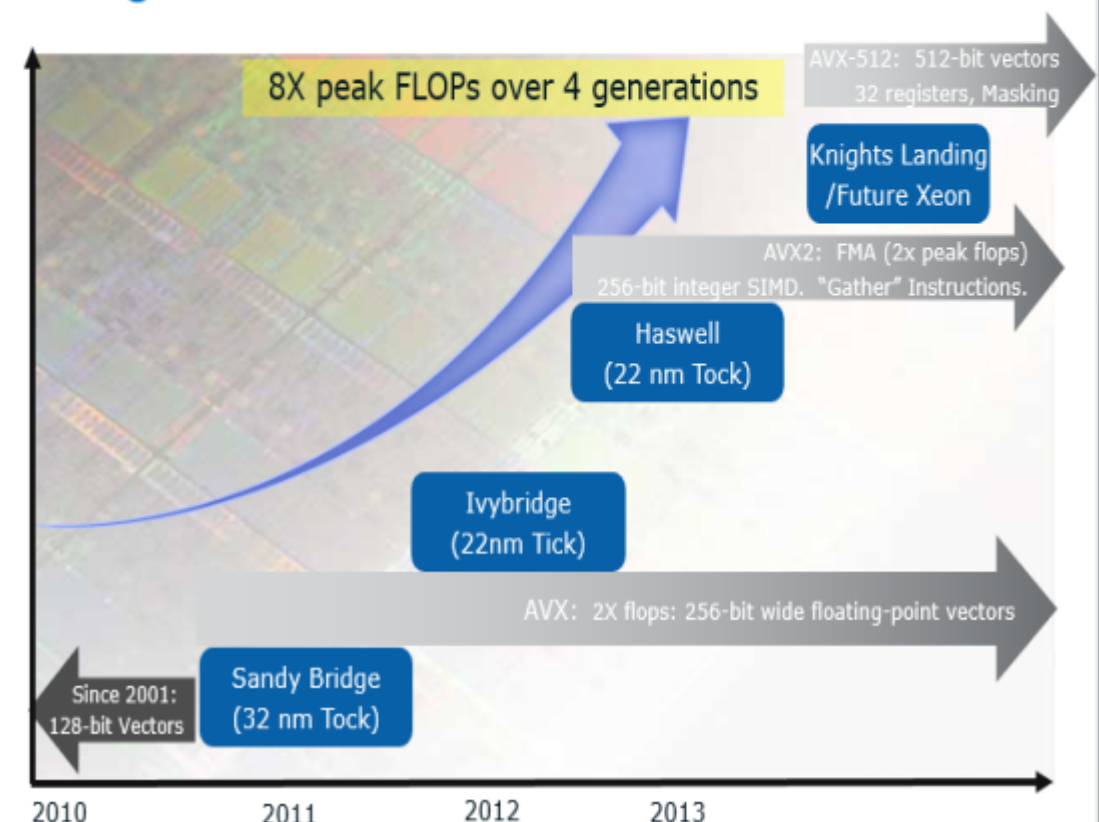
Особенности векторизации кода

Эволюция архитектур

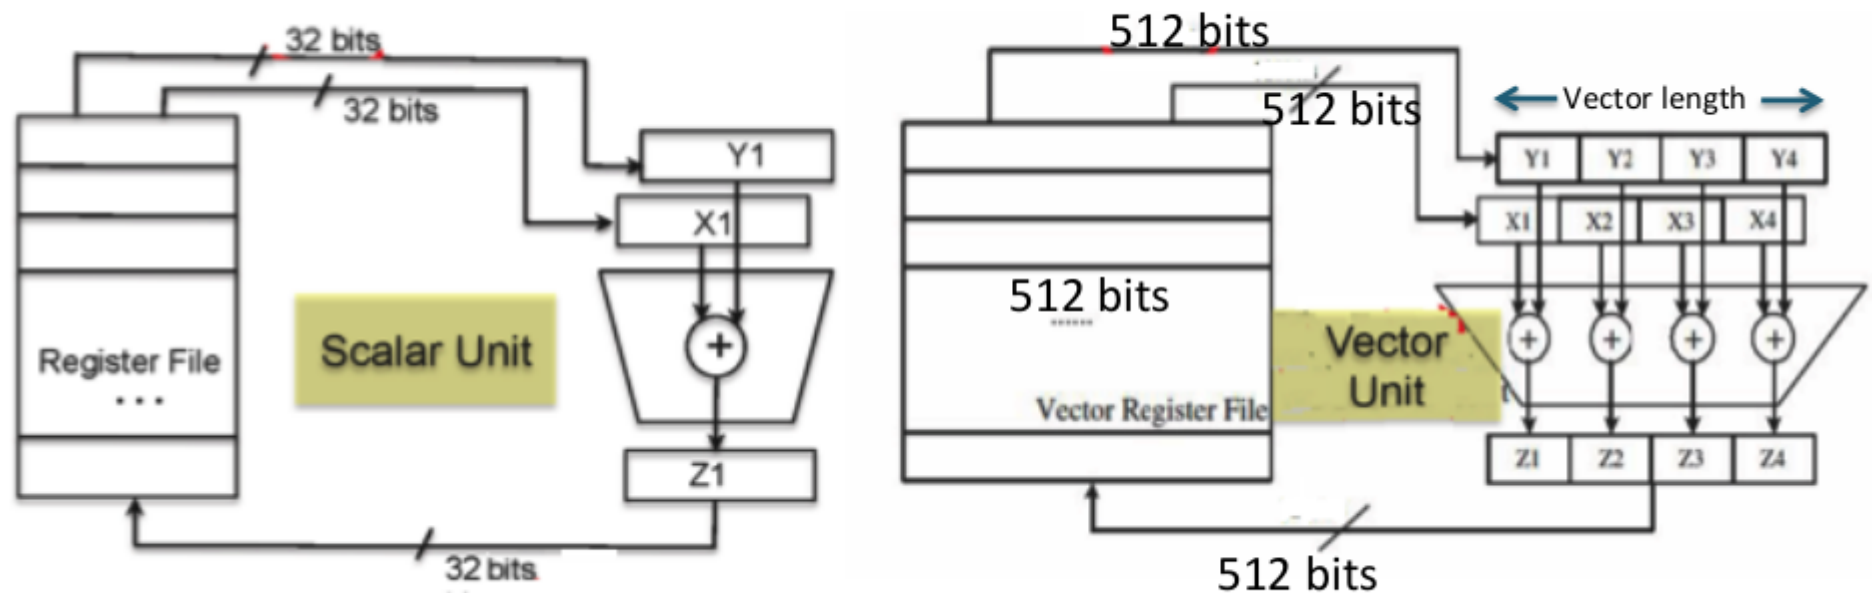


All future processors will have multiple cores with SIMD instructions

Intel® Advanced Vector Extensions



Архитектуры SIMD современных ядер



Intel SIMD ISA Evolution

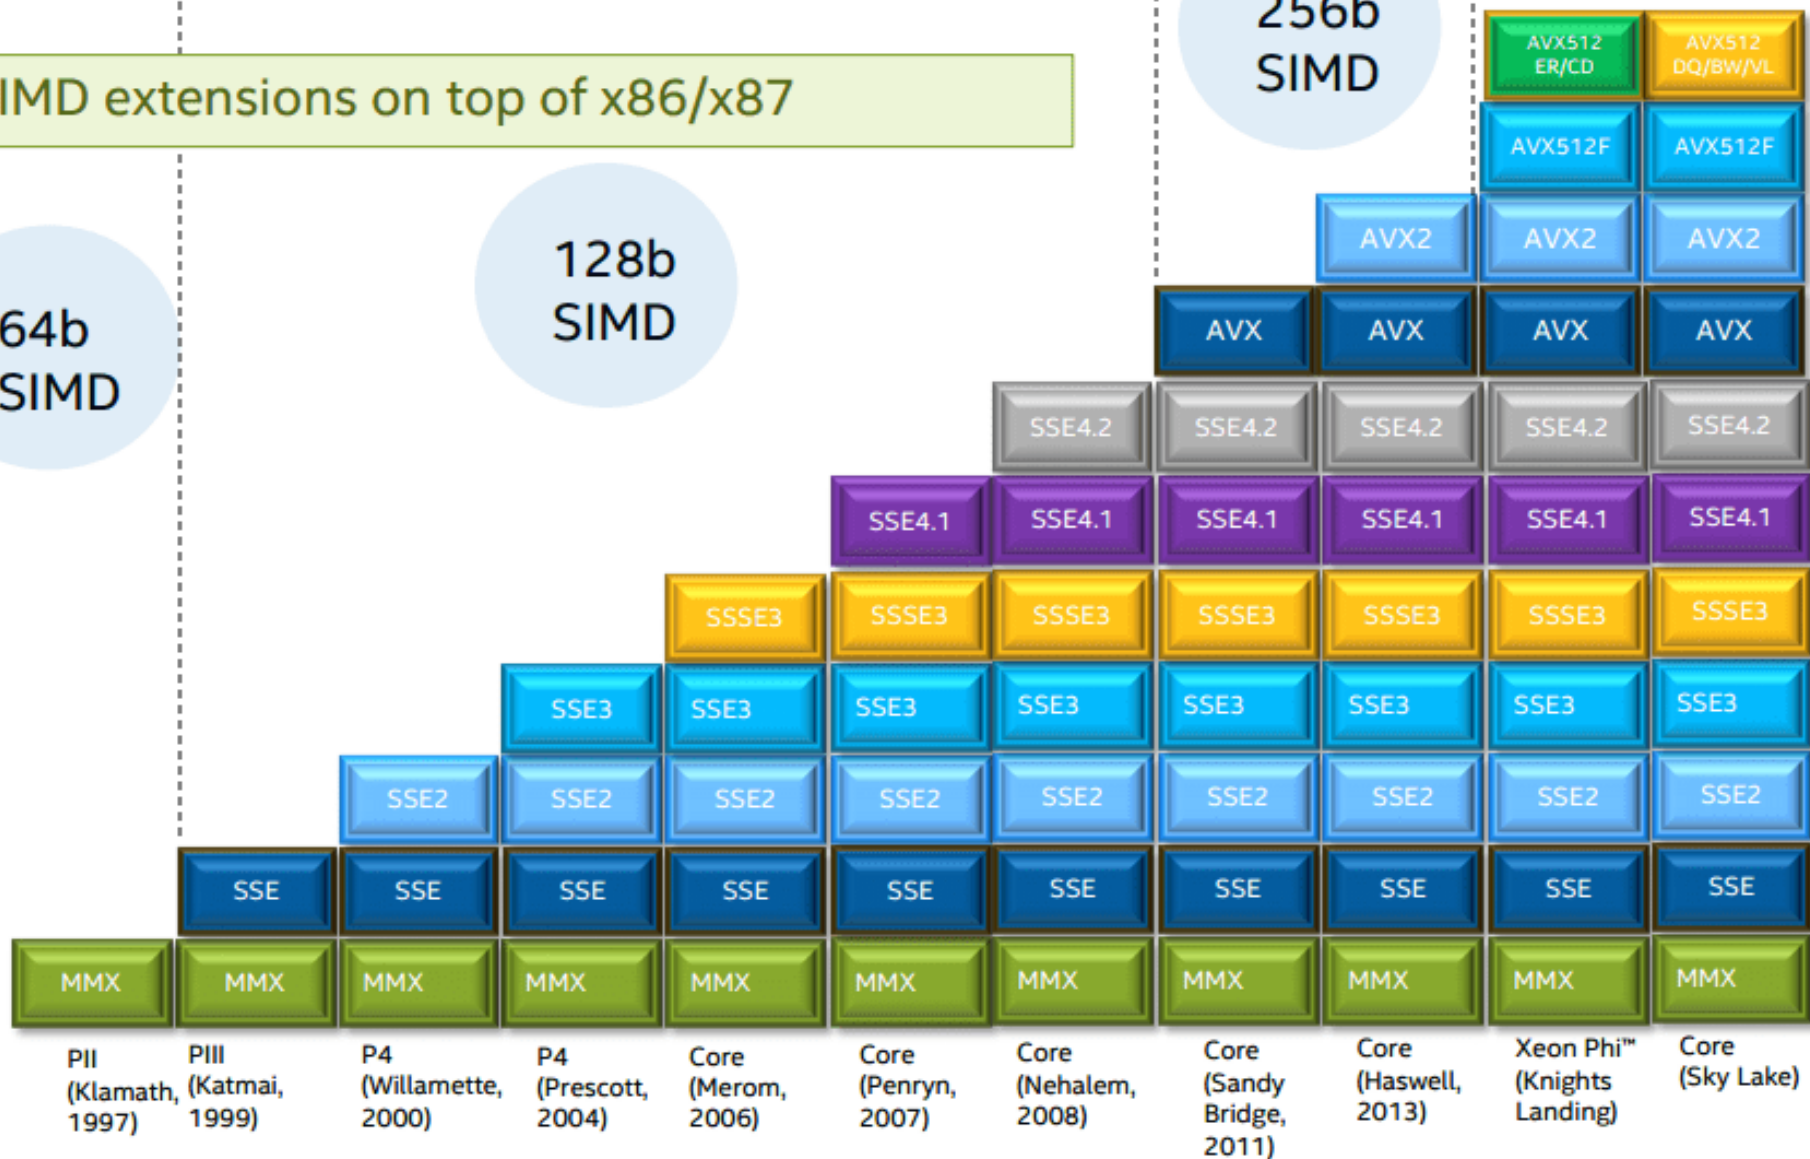
SIMD extensions on top of x86/x87

64b
SIMD

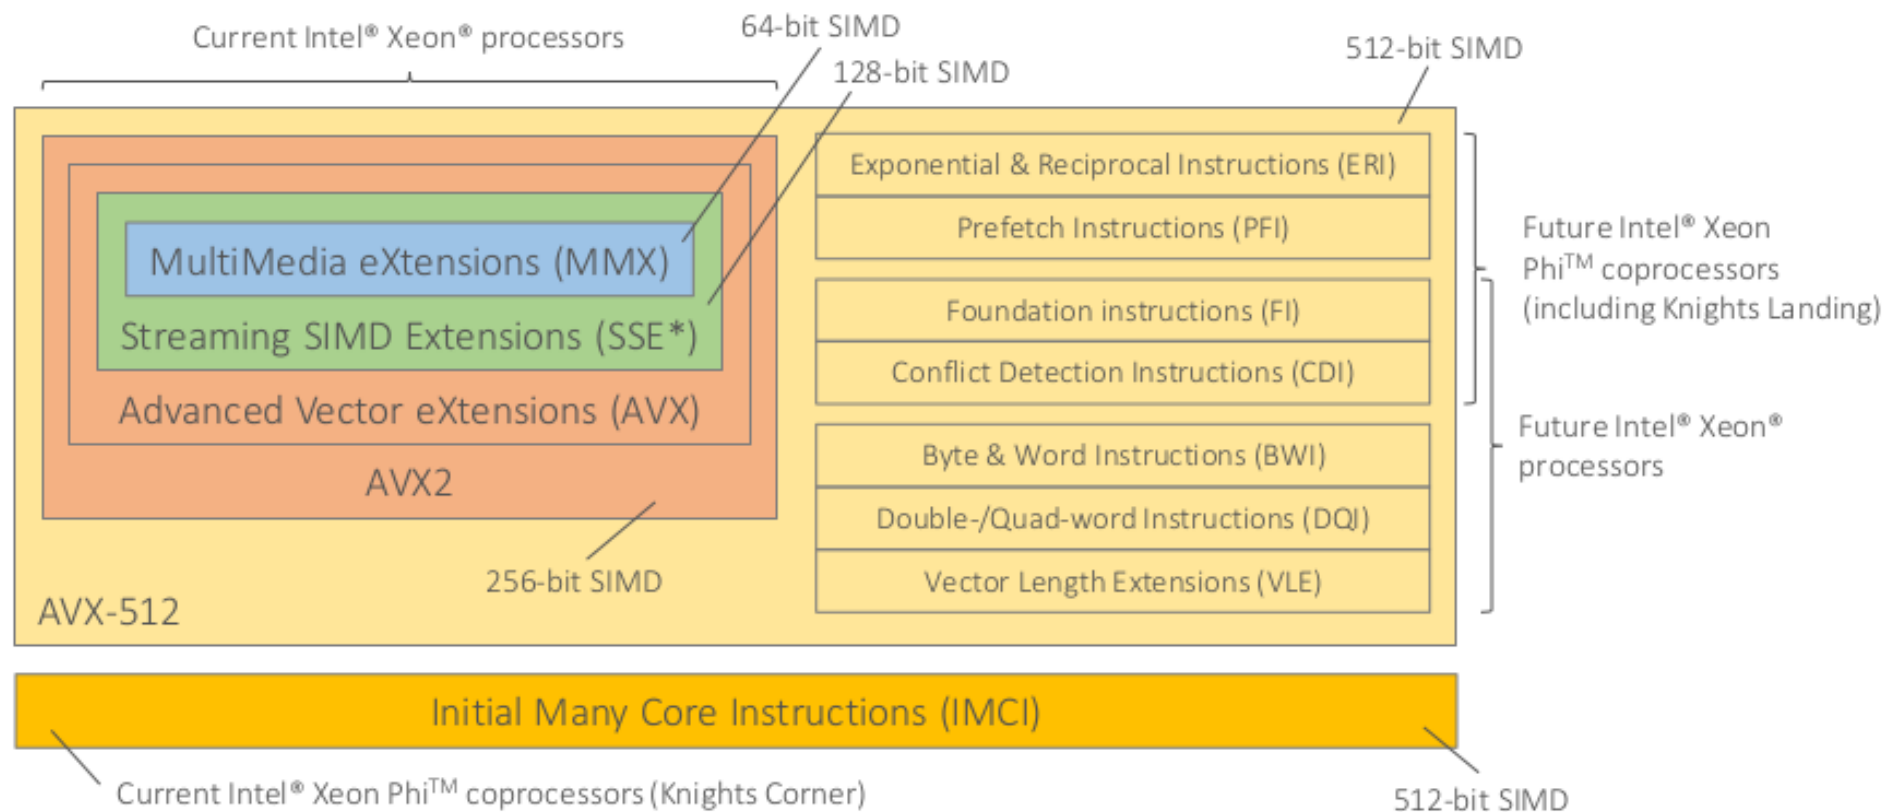
128b
SIMD

256b
SIMD

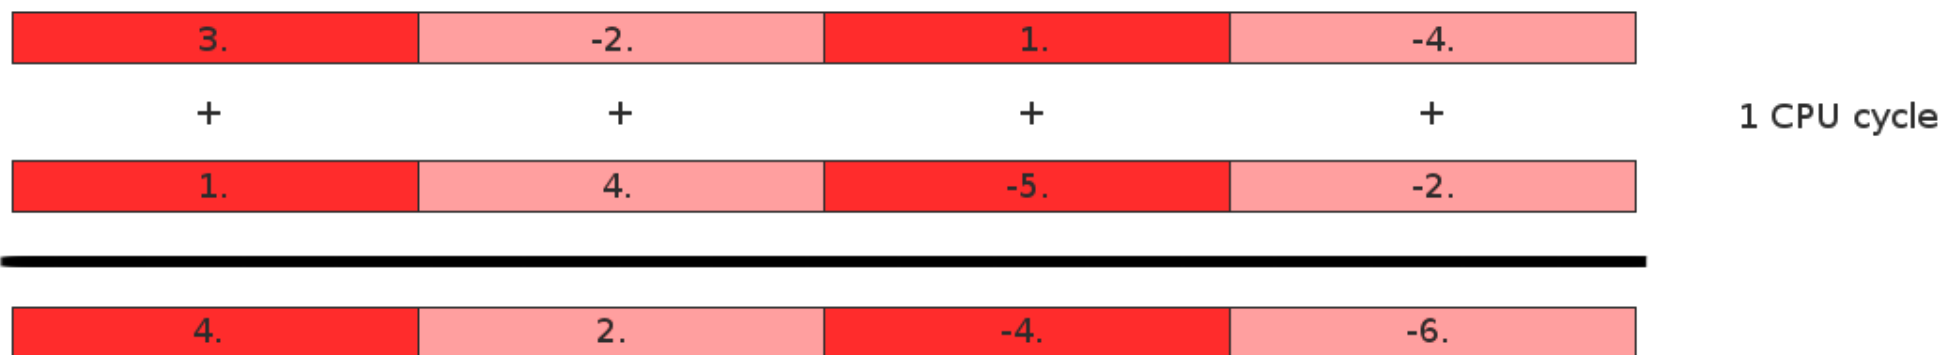
512b
SIMD



Типы инструкций SIMD у решений Intel



Различия для разновидностей типа



Трансформация кода

```
DO I = 1, N  
  Z(I) = X(I) + Y(I)  
ENDDO
```

```
DO I = 1, N, 4  
  Z(I) = X(I) + Y(I)  
  Z(I+1) = X(I+1) + Y(I+1)  
  Z(I+2) = X(I+2) + Y(I+2)  
  Z(I+3) = X(I+3) + Y(I+3)  
ENDDO
```

```
VLOAD  X(I), X(I+1), X(I+2), X(I+3)  
VLOAD  Y(I), Y(I+1), Y(I+2), Y(I+3)  
VADD   Z(I, ..., I+3)  X+Y(I, ..., I+3)  
VSTORE Z(I), Z(I+1), Z(I+2), Z(I+3)
```

Изменение порядка выполнения

```
DO I = 1, 4  
  A(I) = B(I) + C(I)  
  D(I) = E(I) + F(I)  
END DO
```

Non-Vector

```
A(1) = B(1) + C(1)  
  D(1) = E(1) + F(1)  
A(2) = B(2) + C(2)  
  D(2) = E(2) + F(2)  
A(3) = B(3) + C(3)  
  D(3) = E(3) + F(3)  
A(4) = B(4) + C(4)  
  D(4) = E(4) + F(4)
```

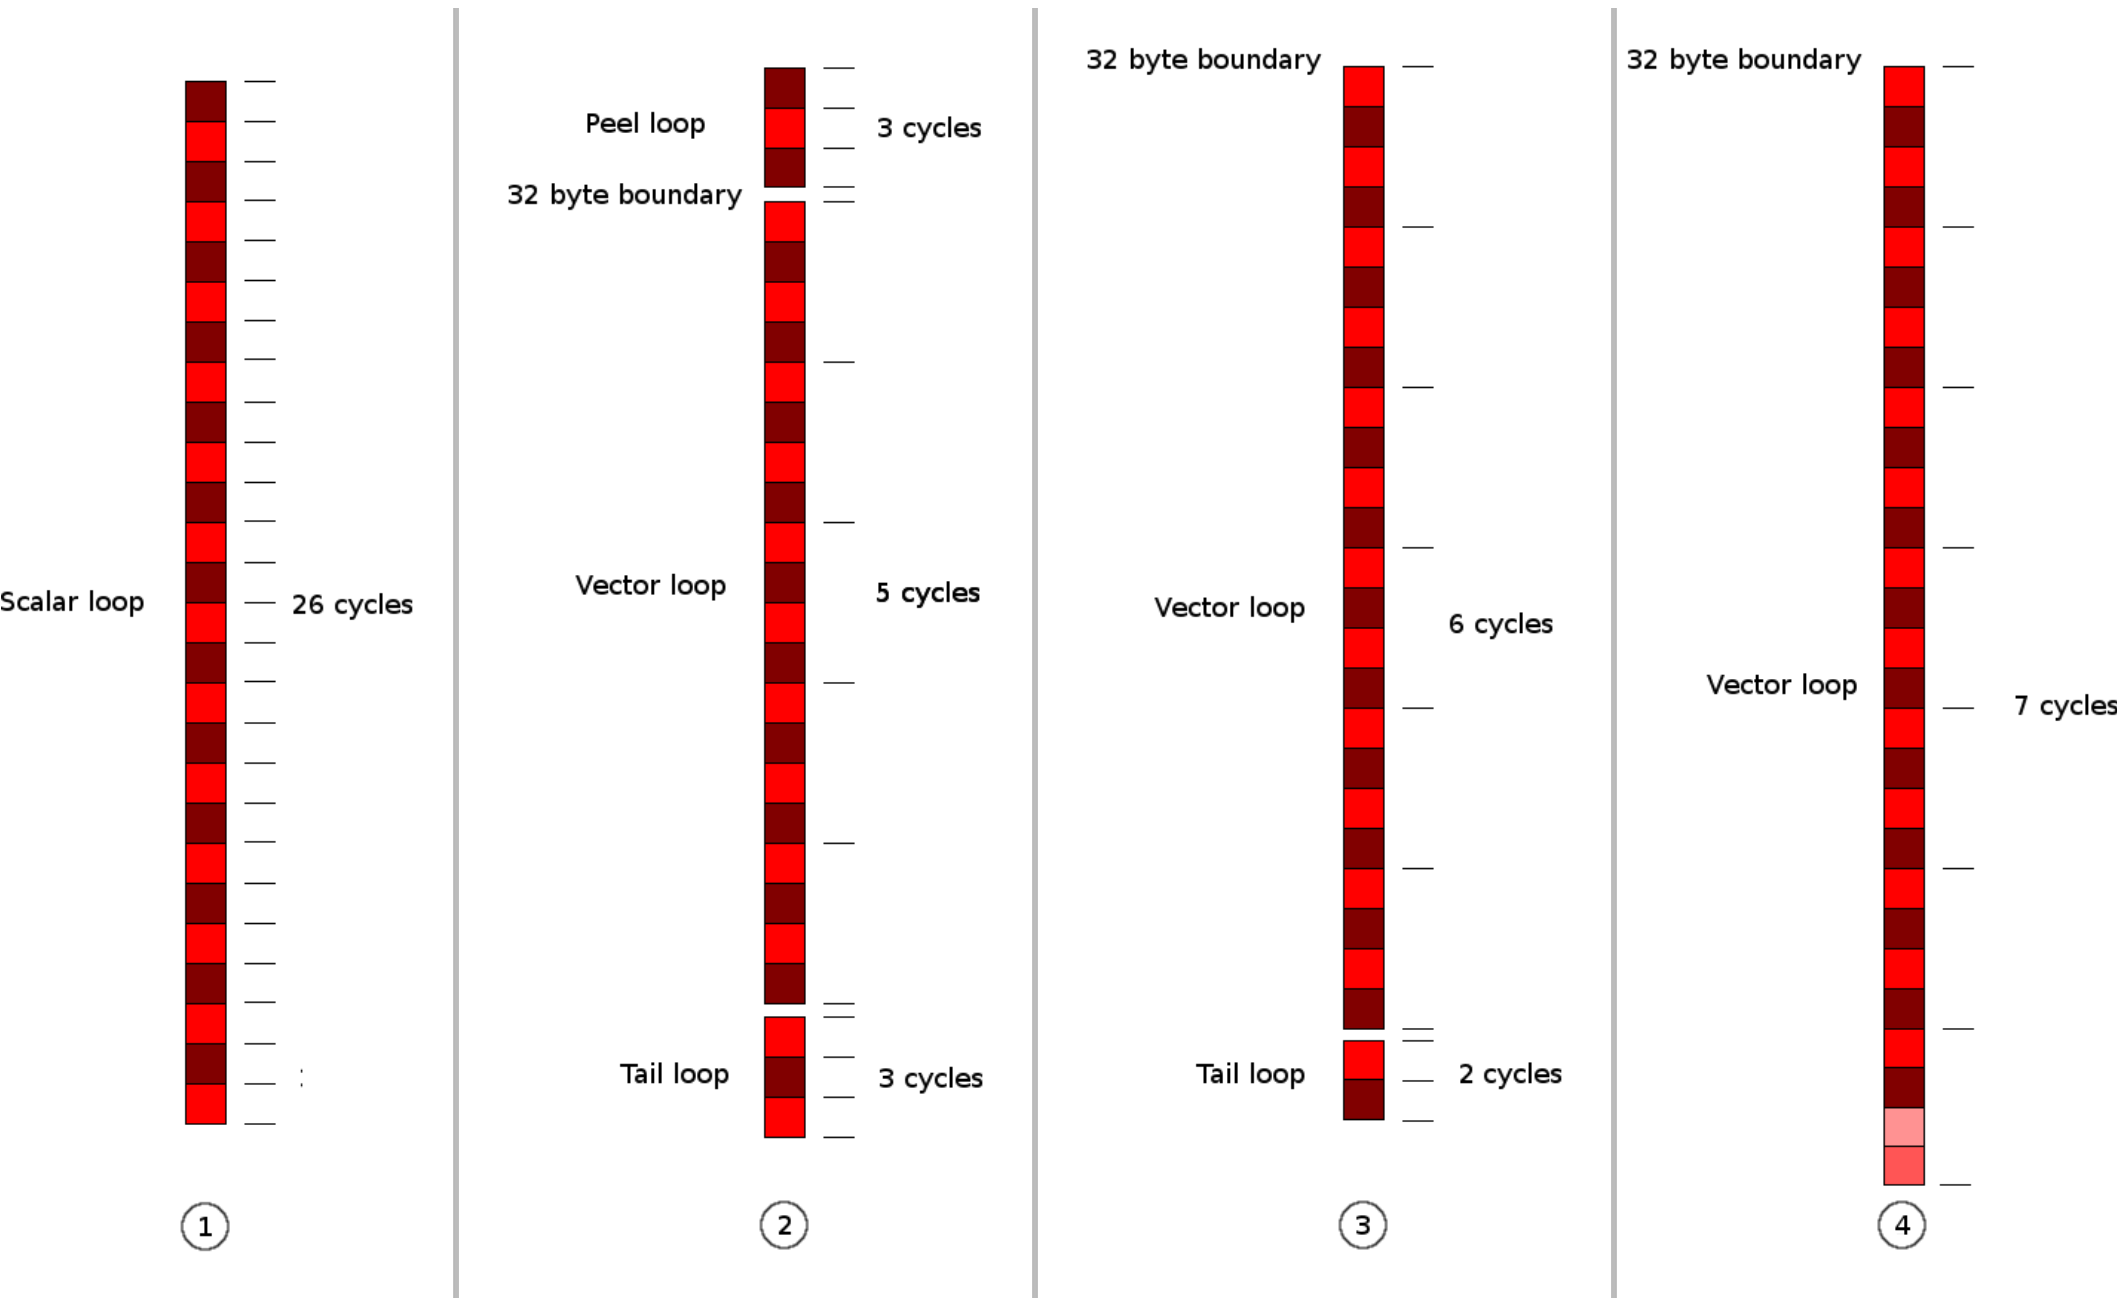
Vector

```
A(1) = B(1) + C(1)  
A(2) = B(2) + C(2)  
A(3) = B(3) + C(3)  
A(4) = B(4) + C(4)  
  D(1) = E(1) + F(1)  
  D(2) = E(2) + F(2)  
  D(3) = E(3) + F(3)  
  D(4) = E(4) + F(4)
```

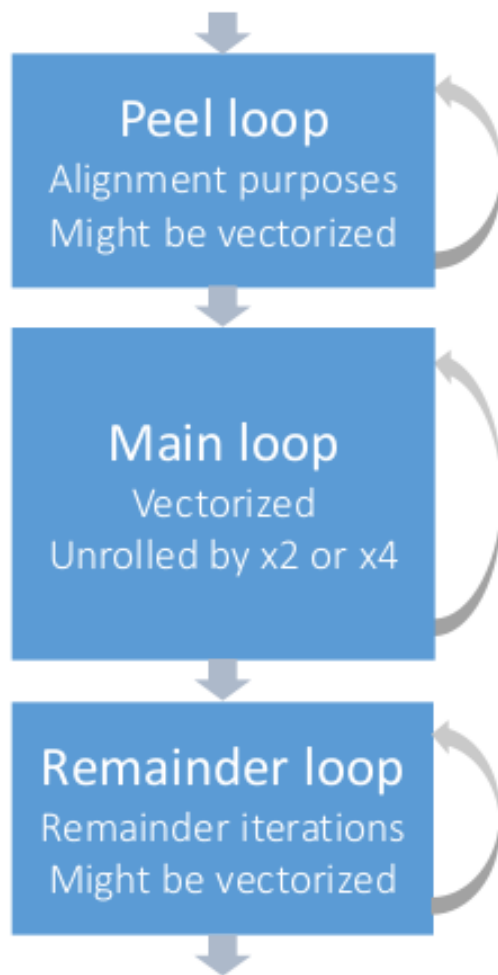
Order of execution



Три части порождаемого цикла



Три части порождаемого цикла



```
LOOP BEGIN at gas_dyn2.f90(2330,26)
<Peeled>
  remark #15389: vectorization support: reference AMAC1U has unaligned access
  remark #15381: vectorization support: unaligned access used inside loop body
  remark #15301: PEEL LOOP WAS VECTORIZED
LOOP END

LOOP BEGIN at gas_dyn2.f90(2330,26)
  remark #25084: Preprocess Loopnests: Moving Out Store
  remark #15388: vectorization support: reference AMAC1U has aligned access
  remark #15399: vectorization support: unroll factor set to 2
  remark #15300: LOOP WAS VECTORIZED
  remark #15475: --- begin vector loop cost summary ---
  remark #15476: scalar loop cost: 8
  remark #15477: vector loop cost: 0.620
  remark #15478: estimated potential speedup: 15.890
  remark #15479: lightweight vector operations: 5
  remark #15488: --- end vector loop cost summary ---
  remark #25018: Total number of lines prefetched=4
  remark #25019: Number of spatial prefetches=4, dist=8
  remark #25021: Number of initial-value prefetches=6
LOOP END

LOOP BEGIN at gas_dyn2.f90(2330,26)
<Remainder>
  remark #15388: vectorization support: reference AMAC1U has aligned access
  remark #15388: vectorization support: reference AMAC1U has aligned access
  remark #15301: REMAINDER LOOP WAS VECTORIZED
LOOP END
```

Условия векторизации

- Данные должны быть сплошными в памяти
- Первый элемент каждого вектора должен быть выровнен
- Не должно быть перекрытий по памяти
- Не должно быть зависимостей итераций в виде чтение-после-записи

На что обращать внимание в коде

```
for (int i = 0; i < N; i++)  
    a[i] = a[i-1] + b[i];
```

```
for (int i = 0; i < N; i++)  
    a[c[i]] = b[d[i]];
```

```
for (int i = 0; i < N; i++)  
    a[i] = foo(b[i]);
```

Для C99 и C++

```
void v_add(float *restrict c,  
           float *restrict a,  
           float *restrict b)  
{  
    for (int i = 0; i < N; i++)  
        c[i] = a[i] + b[i];  
}
```

```
void v_add(float *c, float *a, float *b)  
{  
    #pragma ivdep  
    for (int i = 0; i < N; i++)  
        c[i] = a[i] + b[i];  
}
```

Управляемая векторизация: ЦИКЛЫ

```
void v_add(float *c, float *a, float *b)
{
    #pragma simd
    for (int i = 0; i < N; i++)
        c[i] = a[i] + b[i];
}
```

Управляемая векторизация: OpenMP

```
#pragma omp simd reduction(+:sum) aligned(a : 64)
for(i = 0; i < num; i++) {
    a[i] = b[i] * c[i];
    sum = sum + a[i];
}
```

Управляемая векторизация: функции

```
__declspec(vector)  
void v_add(float c, float a, float b)  
{  
    c = a + b;  
}  
  
...  
for (int i = 0; i < N; i++)  
    v_add(C[i], A[i], B[i]);
```


Управляемая векторизация: функции (Fortran)

```
module fofx
contains
  function f(x) ← Line 7
!dir$ attributes vector :: f      Elemental function in vectorization sense
    real, intent(in) :: x
    real f
    f = cos(x * x + 1.) / (x * x + 1.)
  end function
end module

program main
  use fofx
  real a(100), x(100)
  ...
  do i=1,100
    a(i) = f(x(i))
  end do
  ...
end program
```

```
$ ifort -vec-report=3 elemental.f90
...
elemental.f90(67): (col. 11) remark: LOOP WAS
VECTORIZED
...
elemental.f90(7): (col. 18) remark: FUNCTION WAS
VECTORIZED
elemental.f90(7): (col. 18) remark: FUNCTION WAS
VECTORIZED
Line 67
```

Управляемая векторизация: функции (Fortran)

```
module fofx
contains
  elemental function f(x)
    !dir$ attributes vector :: f
    real, intent(in) :: x
    real f
    f = cos(x * x + 1.) / (x * x + 1.)
  end function
end module

program main
  use fofx
  real a(100), x(100)
  ...
  a = f(x)
  ...
end program
```

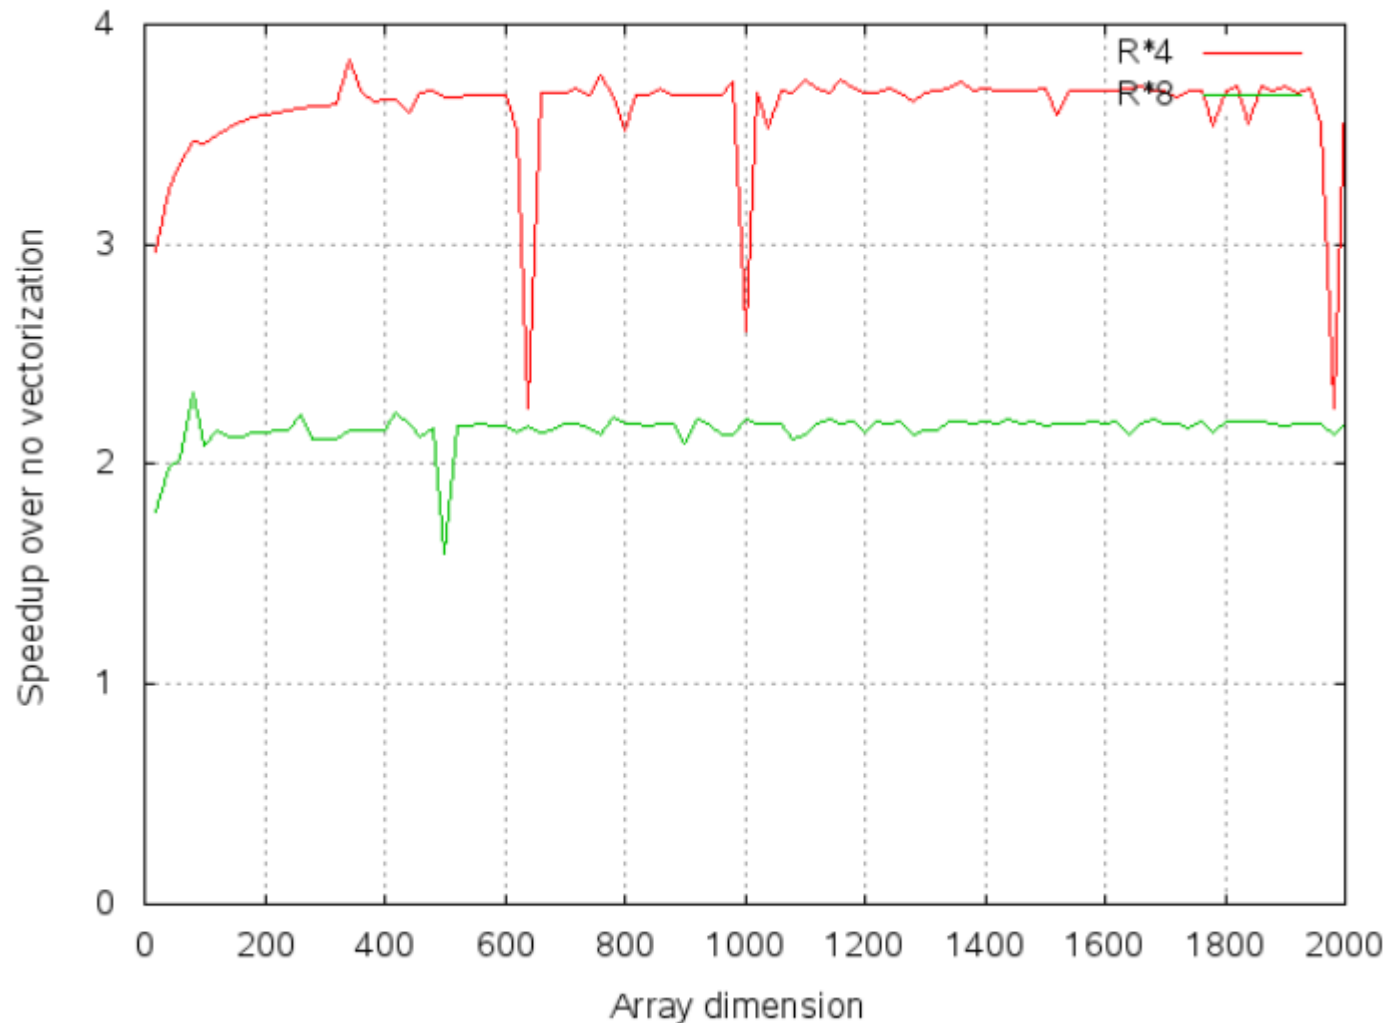
This Fortran 'elemental' clause nothing to do with vect.
Elemental function in vectorization sense

```
$ ifort -vec-report=3 elemental.f90
...
elemental.f90(67): (col. 11) remark: LOOP WAS
VECTORIZED
...
elemental.f90(7): (col. 28) remark: FUNCTION WAS
VECTORIZED
elemental.f90(7): (col. 28) remark: FUNCTION WAS
VECTORIZED
```

Line 67



Прирост производительности для элементных функций



Управляемая векторизация: OpenMP

```
#pragma omp parallel for simd  
for(i = 0; i < num; i++) {  
    sum = sum + a[i];  
}
```

Управляемая векторизация: OpenMP

```
#pragma omp declare simd  
float myfunction(float a, float b, float c )  
{  
    return a * b + c ;  
}
```

```
#pragma omp simd  
for(i = 0; i < num; i++) {  
    OUT[i] = myfunction(arraya[i], arrayb[i], arrayc[i]);  
}
```

Явная векторизация

```
a[:]           // All elements  
a[2:6]         // Elements 2 to 7  
a[:,5]         // Column 5  
a[0:3:2]       // Elements 0,2,4
```

```
__declspec(vector)  
void v_add(float c, float a, float b)  
{  
    c = a + b;  
}  
...  
v_add(C[:, A[:, B[:, ]]);
```

Массивы структур и структуры массивов

```
// Array of Structures (AoS)
struct coordinate {
    float x, y, z;
} crd[N];

...
for (int i = 0; i < N; i++)
    ... = ... f(crd[i].x, crd[i].y, crd[i].z);
```

Consecutive elements in memory →

x0 y0 z0 x1 y1 z1 ... x(n-1) y(n-1) z(n-1)

```
// Structure of Arrays (SoA)
struct coordinate {
    float x[N], y[N], z[N];
} crd;

...
for (int i = 0; i < N; i++)
    ... = ... f(crd.x[i], crd.y[i], crd.z[i]);
```

Consecutive elements in memory →

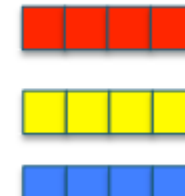
x0 x1 ... x(n-1) y0 y1 ... y(n-1) z0 z1 ... z(n-1)

Массив структур

```
type coords
  real :: x, y, z
end type
type (coords) :: p(100)
real dsquared(100)
```



```
do i=1,100
  dsquared(i) = p(i)%x**2 + p(i)%y**2 + p(i)%z**2
end do
```



Структура массивов

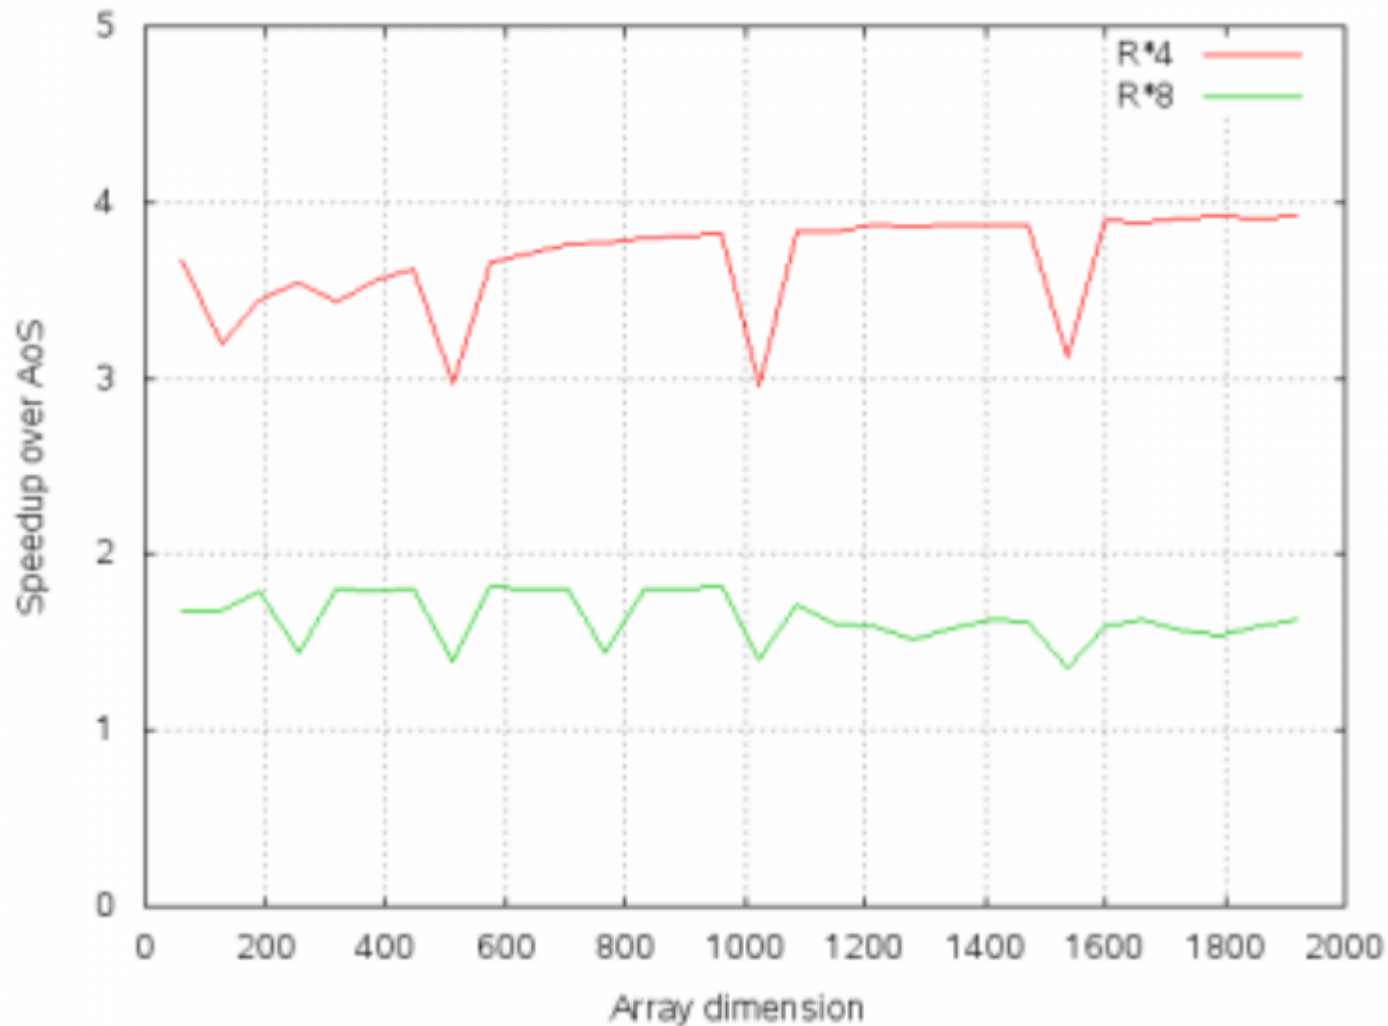
```
type coords
  real :: x(100), y(100), z(100)
end type
type (coords) :: p
real dsquared(100)
```



```
do i=1,100
  dsquared(i) = p%x(i)**2 + p%y(i)**2 + p%z(i)**2
end do
```



Прирост производительности для структуры массивов



Используйте средства компиляторов

```
edison11 h/hjw> cc -vec-report=6 -c mm.c
```

```
mm.c(7): (col. 2) remark: loop was not vectorized: loop was transformed to memset or memcpy
```

```
mm.c(10): (col. 2) remark: vectorization support: reference c has aligned access
```

```
mm.c(10): (col. 2) remark: vectorization support: reference c has aligned access
```

```
mm.c(10): (col. 2) remark: vectorization support: reference a has aligned access
```

```
mm.c(6): (col. 2) remark: vectorization support: unroll factor set to 4
```

```
mm.c(6): (col. 2) remark: PERMUTED LOOP WAS VECTORIZED
```

```
mm.c(9): (col. 2) remark: loop was not vectorized: not inner loop
```

```
mm.c(7): (col. 2) remark: loop was not vectorized: not inner loop
```

Используйте средства компиляторов: Cray

```
57.  + 1-----<          do it=1,itmax
58.  + 1 br4-----<        do j=1,n
59.  + 1 br4 b-----<      do k=1,n
60.    1 br4 b Vr2--<        do i=1,nr
61.    1 br4 b Vr2          c(i,j) = c(i,j) + a(i,k) * b(k,j)
62.    1 br4 b Vr2-->        end do
63.    1 br4 b----->      end do
64.    1 br4----->        end do
65.    1----->          end do
```

b - blocked
r - unrolled
V - Vectorized

ftn-6254 ftn: VECTOR File = matmat.F, Line = 57

A loop starting at line 57 was not vectorized because a recurrence was found on "c" at line 61

ftn-6294 ftn: VECTOR File = matmat.F, Line = 58

A loop starting at line 58 was not vectorized because a better candidate was found at line 60.

ftn-6049 ftn: SCALAR File = matmat.F, Line = 58

A loop starting at line 58 was blocked with block size 8.

ftn-6005 ftn: SCALAR File = matmat.F, Line = 58

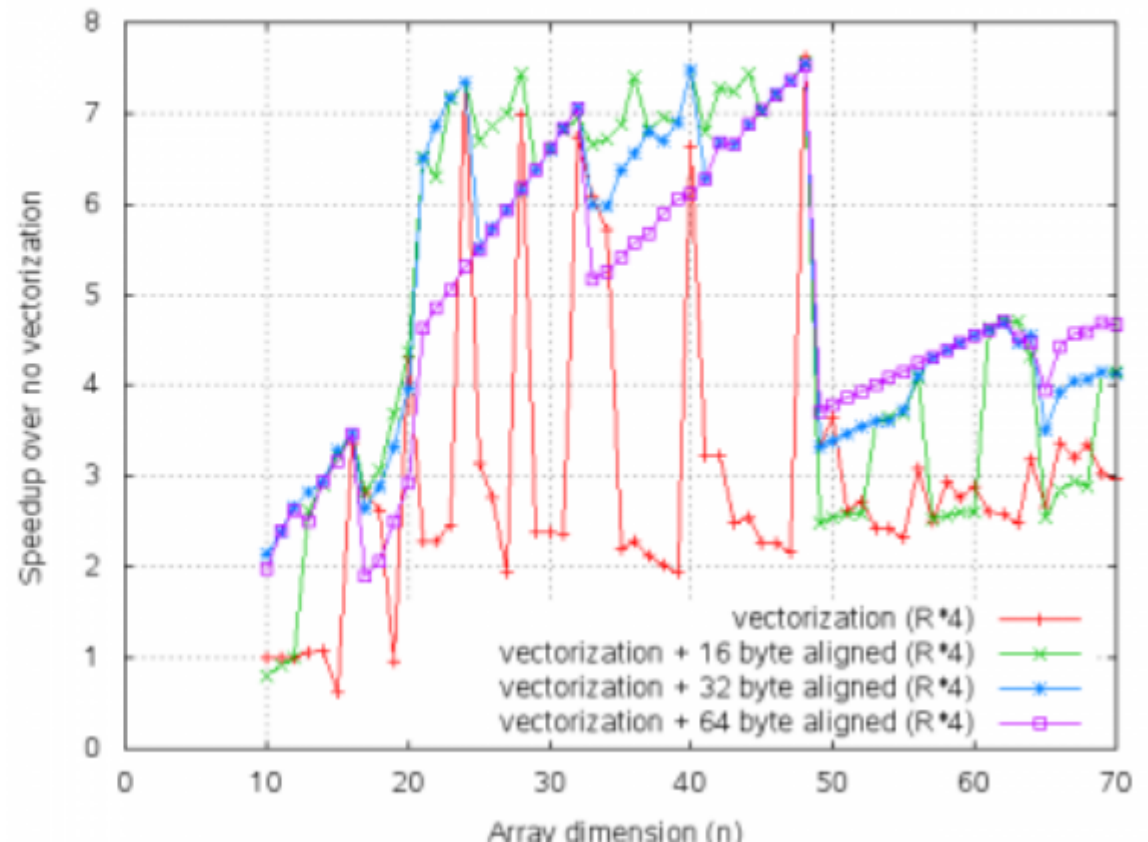
A loop starting at line 58 was unrolled 4 times.

«Наивное» выравнивание данных

```
real, allocatable :: a(:,,:), b(:,,:), c(:,:)  
!dir$ attributes align : 32 :: a,b,c
```

```
...  
allocate (a(npadded,n))  
allocate (b(npadded,n))  
allocate (c(npadded,n))  
...  
do j=1,n  
  do k=1,n  
!dir$ vector aligned  
    do i=1,npadded  
      c(i,j) = c(i,j) &  
        + a(i,k) * b(k,j)  
    end do  
  end do  
end do
```

```
!... Ignore c(n+1:npadded,:)
```



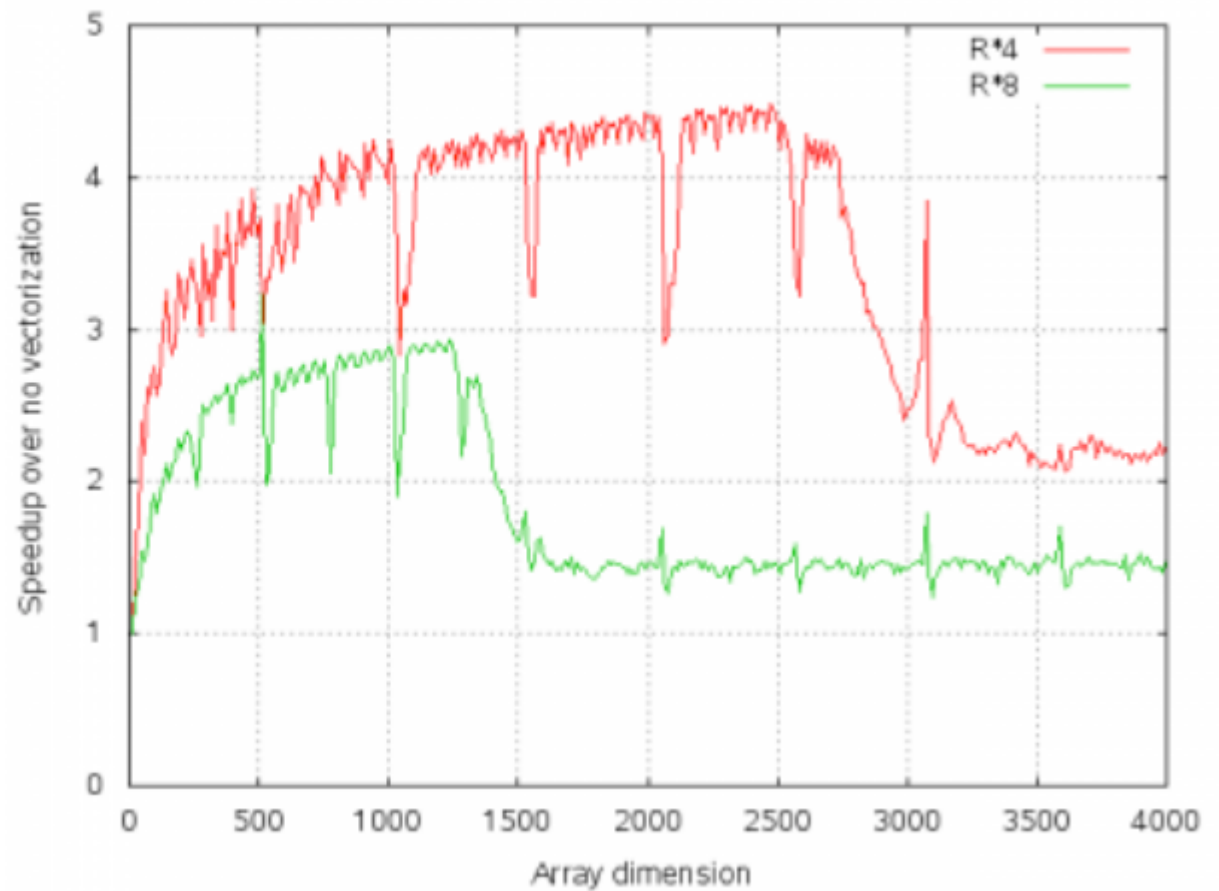
Выравнивание каждой строки

```
% cat Driver.c
...
#define COLBUF 1
...
#define COLWIDTH COL+COLBUF
...
    FTYPE a[ROW][COLWIDTH]    __attribute__((aligned(16)));
    FTYPE b[ROW]               __attribute__((aligned(16)));
    FTYPE x[COLWIDTH]          __attribute__((aligned(16)));
...

% cat Multiply.c
...
#pragma vector aligned
...
    for (j = 0; j < size2; j++) {
        b[i] += a[i][j] * x[j];
    }
```

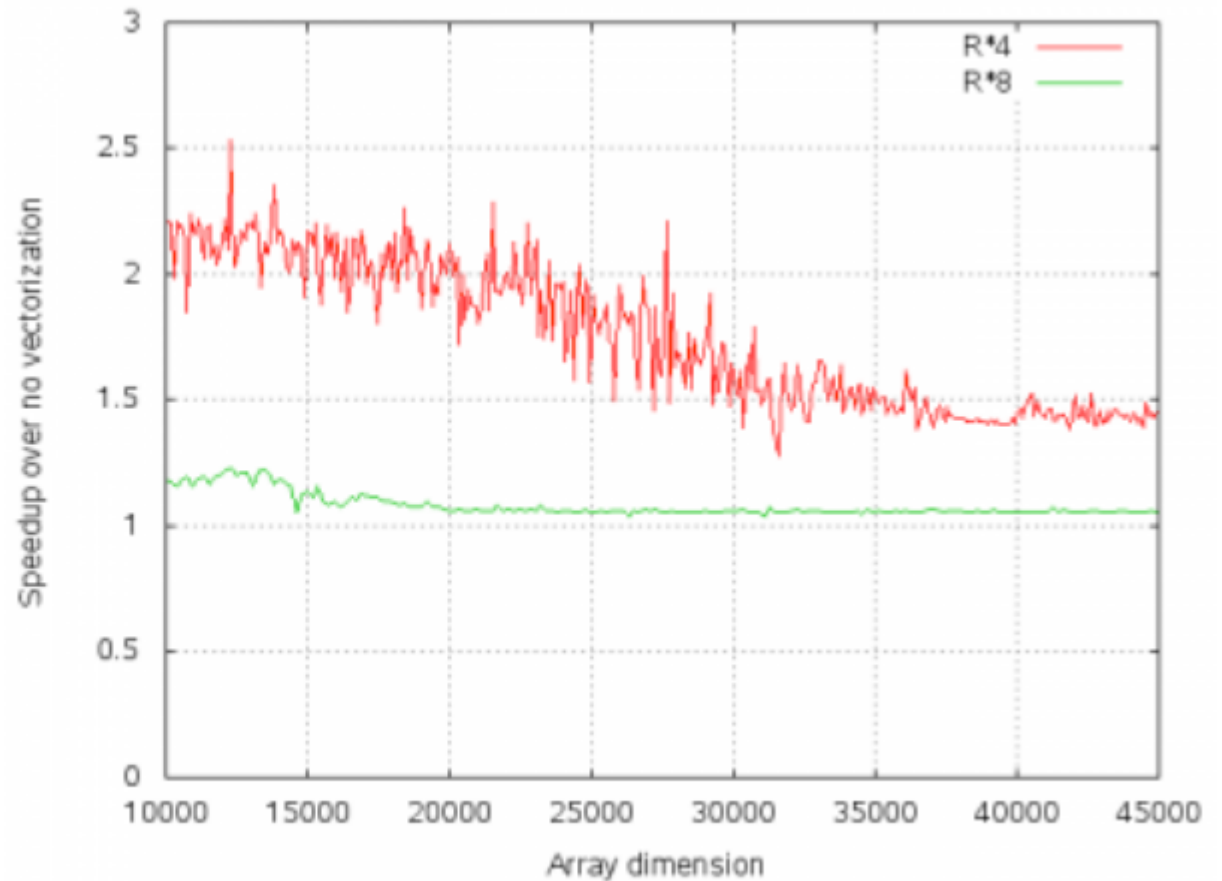
Что ожидать от производительности?

```
do i=1,n  
  c(i) = a(i) + b(i)  
end do
```



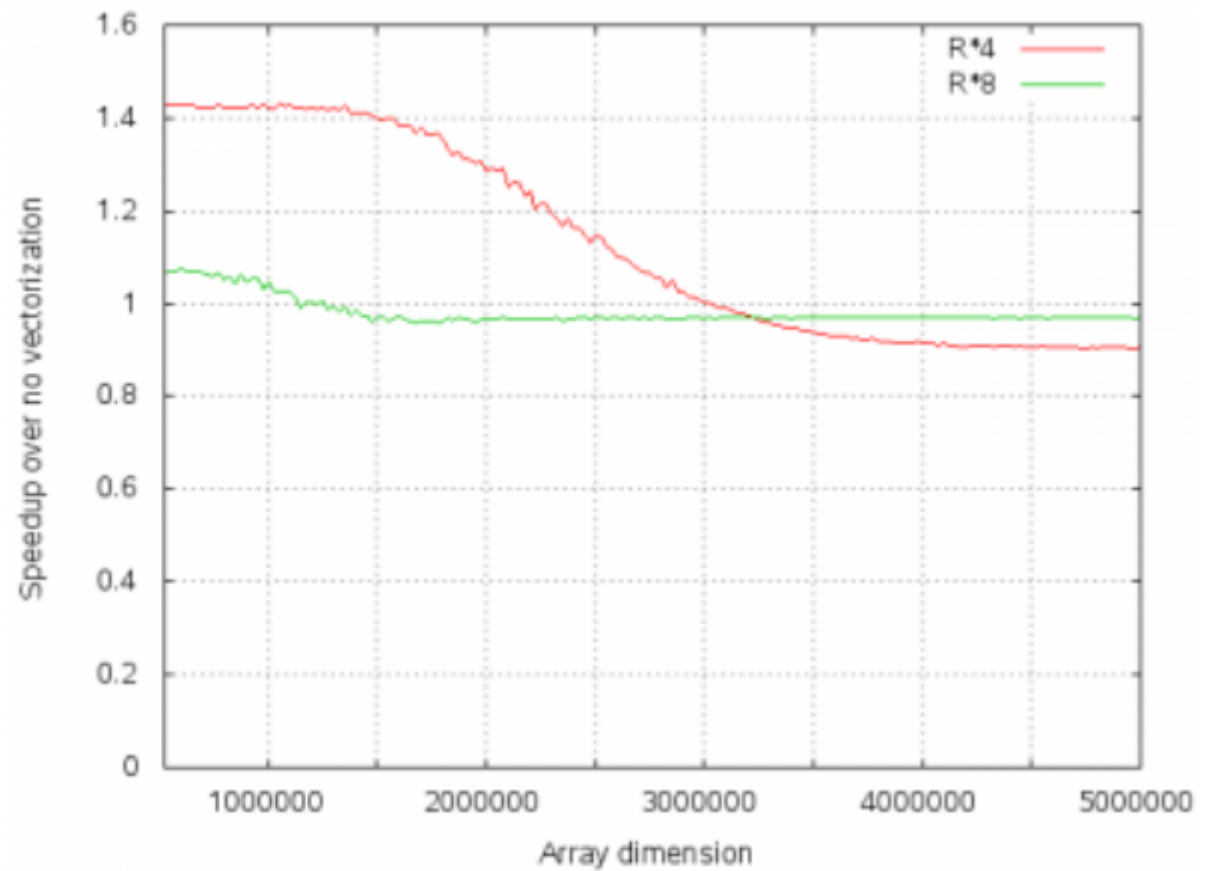
А дальше?

```
do i=1,n  
  c(i) = a(i) + b(i)  
end do
```




?

```
do i=1,n  
  c(i) = a(i) + b(i)  
end do
```



Блочная работа с СОЗУ

```
do i = 1, n
  do j = 1, m
    c += a(i) * b(j)
  enddo
enddo
```



Loads From DRAM:

$$n*m + n$$

```
do jout = 1, m, block
  do i = 1, n
    do j = jout, jout+block-1
      c += a(i) * b(j)
    enddo
  enddo
enddo
```

Loads From DRAM:

$$m/block * (n+block)$$

$$= m*n/block + m$$
